

 VOCABULARY GUIDE

**LOOKS SIMILAR.
MEANS DIFFERENT.**

AI Terms That Confuse Everyone

25 AI terms you've been using interchangeably (but
shouldn't)

→ Swipe to fix your vocabulary



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Latency **VS** Response Time

? WHAT

Latency: time until the first token/byte arrives.

Response Time: total time until the entire response is complete.

✂ WHY THE CONFUSION

Both feel like "how fast is it" but one measures a starting point, the other measures a full journey.

🕒 WHEN IT MATTERS

Chat UIs care about latency (feels instant). Batch jobs care about response time (total throughput).

💡 EXAMPLE

A model with 200ms latency but a 10-second response time feels fast (streaming starts quick) even though the full answer takes long.

✖ COMMON MISTAKE

Teams optimize response time and wonder why users still say the app "feels slow" they ignored latency.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Context Window **vs** Memory

? WHAT

Context Window: the token limit the model can "see" in one session.

Memory: information that persists across sessions.

⌘ WHY THE CONFUSION

Both relate to "what the model remembers" but context window resets, memory doesn't.

🕒 WHEN IT MATTERS

Building a chatbot? Context window decides conversation length. Memory decides if it remembers you tomorrow.

💡 EXAMPLE

ChatGPT's memory feature isn't expanding the context window it's injecting saved facts back into a fresh context window each time.

✖ COMMON MISTAKE

Assuming a bigger context window = the model "remembers" you. It doesn't, once the session ends.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

RAG vs Fine-Tuning

? WHAT

RAG: retrieves external knowledge at query time, no weight changes.

Fine-tuning: retrains the model's weights on new data.

⚡ WHY THE CONFUSION

Both are "ways to make the model know more."

🕒 WHEN IT MATTERS

RAG → fast-changing, factual, document-heavy knowledge.

Fine-tuning → changing tone, behavior, format, or domain reasoning style.

💡 EXAMPLE

A legal chatbot needing today's case law → RAG. A chatbot needing to "sound like a senior partner" → fine-tuning.

✖ COMMON MISTAKE

Fine-tuning a model to "add facts" facts go stale; RAG handles that far better.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Agent **vs** Workflow

? WHAT

Agent: decides its own next steps using reasoning.

Workflow: follows a fixed, predefined sequence.

⚡ WHY THE CONFUSION

Both "automate a process."

🕒 WHEN IT MATTERS

Predictable, repeatable task → workflow (cheaper, reliable). Open-ended, judgment-based task → agent.

💡 EXAMPLE

"Send invoice every Friday" = workflow. "Investigate why revenue dropped and decide what to check next" = agent.

✖ COMMON MISTAKE

Calling every LLM-powered automation an "agent." Most are just workflows with an LLM step.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Embedding **VS** Vector

? WHAT

Embedding: the process/technique of converting data into numeric representation.

Vector: the output the actual array of numbers.

⚡ WHY THE CONFUSION

People use them as synonyms in casual speech.

🕒 WHEN IT MATTERS

Talking architecture: "we embed the text" (verb/process) vs "we store the vector" (noun/result).

💡 EXAMPLE

"The embedding model converts text into a 1536-dimension vector" one sentence, both terms, both correct.

✖ COMMON MISTAKE

Saying "vector model" instead of "embedding model" the model embeds; the vector is what it produces.



AI Coach John

Follow



 Repost

MCP **VS** API

🔍 WHAT

API: a general interface contract for two systems to talk.

MCP (Model Context Protocol): a standardized protocol specifically for exposing tools/context to LLMs.

⚡ WHY THE CONFUSION

MCP is technically built on API-like communication, so people think it's "just another API."

🕒 WHEN IT MATTERS

Building one-off integration → API. Building something reusable across multiple AI agents/clients → MCP.

💡 EXAMPLE

A weather API returns JSON for any app. An MCP server exposes that same data in a way any LLM client (Claude, etc.) can discover and use automatically.

❌ COMMON MISTAKE

Thinking MCP replaces APIs. It doesn't MCP often wraps APIs to make them LLM-consumable.



Precision **VS** Recall

? WHAT

Precision: of what you retrieved, how much was relevant.

Recall: of all relevant items that exist, how much did you retrieve.

✂ WHY THE CONFUSION

Both sound like "accuracy" of a retrieval system.

🕒 WHEN IT MATTERS

Search engines often favor precision (top results matter).
Legal/medical search favors recall (can't miss anything).

💡 EXAMPLE

Retrieving 3 docs, all relevant = high precision. But if 10 relevant docs existed, recall is low.

✖ COMMON MISTAKE

Optimizing only for precision in RAG systems leads to confident answers built on incomplete context.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Tool Calling **VS** Function Calling

? WHAT

Function Calling: OpenAI's specific term/implementation for invoking a single defined function.

Tool Calling: the broader concept functions, APIs, plugins, retrieval, code execution, etc.

✂ WHY THE CONFUSION

Function calling was the first popularized version, so it became the generic term.

🕒 WHEN IT MATTERS

Cross-provider work Claude, Gemini, and others use "tool calling" as the umbrella term; "function calling" is narrower and OpenAI-specific.

💡 EXAMPLE

A "tool" can be a calculator function, a web search, or a database query function calling usually refers to just the function-style invocation.

✖ COMMON MISTAKE

Using the terms interchangeably in technical docs confuses readers across different model providers.



AI Coach John

Follow



 Repost

GPT vs ChatGPT

? WHAT

GPT: the underlying model/architecture (e.g., GPT-4).

ChatGPT: the product the chat interface, memory, plugins, and UX wrapped around the model.

✂ WHY THE CONFUSION

Most people only ever interact with ChatGPT, so they think "GPT" and "ChatGPT" are the same thing.

🕒 WHEN IT MATTERS

Developers use GPT via API (no chat UI, no memory by default). End users use ChatGPT (full product experience).

💡 EXAMPLE

You can use GPT-4 through the API with zero chat history that's not "using ChatGPT."

✖ COMMON MISTAKE

Saying "ChatGPT's API" there is no such thing; it's the GPT model's API.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

🔄 Repost

Vector Database **vs** Knowledge Graph

? WHAT

Vector Database: stores embeddings, retrieves by similarity.

Knowledge Graph: stores entities and explicit relationships between them.

⌘ WHY THE CONFUSION

Both are "ways to store knowledge for AI to query."

🕒 WHEN IT MATTERS

Fuzzy, semantic similarity search → vector DB. Precise relationship reasoning ("who reports to whom") → knowledge graph.

💡 EXAMPLE

"Find documents similar to this one" → vector DB. "Find all employees who report to someone in Finance" → knowledge graph.

✖ COMMON MISTAKE

Using a vector DB for relationship-heavy queries it can't reliably answer multi-hop relational questions.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

↻ Repost

Token **VS** Word

? WHAT

Token: the model's basic unit of text (subword chunks).

Word: a human linguistic unit.

⚡ WHY THE CONFUSION

Often roughly 1:1 for simple English words, so people assume they're equivalent.

🕒 WHEN IT MATTERS

Pricing, context limits, and rate limits are all token-based, not word-based.

💡 EXAMPLE

"Tokenization" might split into "Token" + "ization" one word, two tokens.

✖ COMMON MISTAKE

Estimating costs by word count token count (often ~1.3x word count) is what actually bills you.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Training **VS** Inference

? WHAT

Training: the process of teaching the model by adjusting weights.

Inference: the process of using the trained model to generate output.

⌘ WHY THE CONFUSION

Both involve "the model doing something with data."

🕒 WHEN IT MATTERS

Training happens once (or periodically); inference happens every single time you send a prompt.

💡 EXAMPLE

GPT-4 was trained over months on massive datasets. Every ChatGPT response you get is inference no learning happens.

✖ COMMON MISTAKE

Thinking the model "learns" from your conversation. It doesn't that's inference, not training.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Fine-tuning **VS** Prompt Engineering

? WHAT

Fine-tuning: updates the model's internal weights with new examples.

Prompt Engineering: crafts better inputs without touching the model at all.

⚡ WHY THE CONFUSION

Both aim to "get better outputs."

🕒 WHEN IT MATTERS

Need quick iteration, low cost → prompt engineering. Need consistent behavior baked in at scale → fine-tuning.

💡 EXAMPLE

Adding "respond only in bullet points" to your prompt = prompt engineering. Training the model on 10,000 bullet-formatted examples = fine-tuning.

✖ COMMON MISTAKE

Jumping to fine-tuning before exhausting prompt engineering most problems are solvable with better prompts, cheaper and faster.



Similarity Search **VS** Semantic Search

? WHAT

Similarity Search: pure vector-distance matching (cosine similarity, etc.).

Semantic Search: broader system that incorporates meaning/intent, often combining similarity search with ranking, filters, or hybrid methods.

⌘ WHY THE CONFUSION

Semantic search is often implemented using similarity search, so the terms blur.

🕒 WHEN IT MATTERS

Designing a search system similarity search is one component; semantic search is the user-facing capability.

💡 EXAMPLE

Pure similarity search might return "Apple the fruit" for a query about "Apple the company" if vectors are close. Semantic search systems add context/ranking to avoid that.

✖ COMMON MISTAKE

Calling raw cosine-similarity retrieval "semantic search" it's a building block, not the whole system.



AI Coach John

Follow



 Repost

Dense Retrieval **VS** Sparse Retrieval

? WHAT

Dense Retrieval: uses embeddings to find conceptually similar text.

Sparse Retrieval: uses keyword/term matching (e.g., BM25, TF-IDF).

⚡ WHY THE CONFUSION

Both are "retrieval methods" used in RAG pipelines.

🕒 WHEN IT MATTERS

Conceptual/paraphrased queries → dense. Exact keyword/legal/code search → sparse. Best systems use both (hybrid retrieval).

💡 EXAMPLE

Searching "car problems" finding a doc about "automobile issues" = dense retrieval working. Searching an exact product SKU = sparse retrieval winning.

❌ COMMON MISTAKE

Going all-in on dense retrieval and losing exact-match accuracy for codes, IDs, or names.



AI Coach John

Follow



 Repost

Reflection **VS** Self-Correction

? WHAT

Reflection: the model evaluates/critiques its own output.

Self-Correction: the model acts on that critique to produce a better output.

⌘ WHY THE CONFUSION

They're sequential steps of the same loop, so people merge them into one idea.

🕒 WHEN IT MATTERS

Designing agent loops reflection without self-correction is just commentary; you need both for actual improvement.

💡 EXAMPLE

"This answer is missing a citation" = reflection. Regenerating the answer with the citation added = self-correction.

✖ COMMON MISTAKE

Building a "reflection step" that never feeds back into the generation wasted compute, no improvement.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Autonomous **VS** Automated

? WHAT

Automated: follows fixed, predefined rules no decision-making.

Autonomous: makes independent decisions based on context/reasoning.

⌘ WHY THE CONFUSION

Both mean "runs without a human."

🕒 WHEN IT MATTERS

Marketing vs. reality many "autonomous agents" are actually just automated scripts with an LLM bolted on.

💡 EXAMPLE

A cron job emailing a report = automated. An agent that decides whether a report is even needed today = autonomous.

✖ COMMON MISTAKE

Overselling automation as autonomy sets wrong expectations for what the system can actually handle.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

↻ Repost

Hallucination VS Bias

? WHAT

Hallucination: the model generates false/fabricated information confidently.

Bias: the model's outputs are systematically skewed toward certain patterns, viewpoints, or groups.

⚡ WHY THE CONFUSION

Both are "model gets it wrong" but one is about truth, the other about fairness/skew.

🕒 WHEN IT MATTERS

Debugging output errors fact-check for hallucination, audit for bias; different fixes for each.

💡 EXAMPLE

Model citing a fake research paper = hallucination. Model consistently associating certain job titles with one gender = bias.

✖ COMMON MISTAKE

Lumping all "wrong" outputs under "hallucination" bias often needs dataset-level fixes, not fact-checking.



Accuracy VS Precision

? WHAT

Accuracy: overall correctness across all predictions (right vs total).

Precision: correctness specifically among positive predictions.

⚡ WHY THE CONFUSION

Both sound like "how often is the model right."

🕒 WHEN IT MATTERS

Imbalanced datasets accuracy can be misleading; precision tells you a sharper story.

💡 EXAMPLE

A spam filter that marks everything "not spam" can have 95% accuracy (if spam is rare) but 0% precision on catching spam.

✖ COMMON MISTAKE

Reporting only accuracy on imbalanced classification tasks it hides real performance issues.



AI Coach John

Follow



 Repost

Monitoring **VS** Observability

? WHAT

Monitoring: tracking known, predefined metrics (latency, error rate, token usage).

Observability: the ability to debug unknown issues using logs, traces, and system internals.

⚡ WHY THE CONFUSION

Both are "watching the system."

🕒 WHEN IT MATTERS

Monitoring tells you something broke. Observability tells you why it broke.

💡 EXAMPLE

Monitoring shows latency spiked at 2pm. Observability lets you trace it to a specific slow tool call inside an agent chain.

✖ COMMON MISTAKE

Having dashboards (monitoring) but no tracing (observability) you see the symptom, never the cause.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Open Source **VS** Open Weight

? WHAT

Open Source: full transparency code, training data, methodology all available.

Open Weight: only the trained model parameters are released; training data/code stay private.

⚡ WHY THE CONFUSION

Both get marketed as "open."

🕒 WHEN IT MATTERS

Compliance, auditing, and reproducibility open weight models can't be fully audited or retrained from scratch.

💡 EXAMPLE

Llama models are open weight (you get the weights, not the full training pipeline or dataset).

✖ COMMON MISTAKE

Calling open-weight models "open source" a common but technically inaccurate label.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Quantization **VS** Compression

? WHAT

Quantization: reduces numeric precision of weights (e.g., 32-bit → 4-bit).

Compression: the broader umbrella includes quantization, pruning, distillation, and more.

⚡ WHY THE CONFUSION

Quantization is the most common compression technique, so people use it as a stand-in for the whole category.

🕒 WHEN IT MATTERS

Choosing a model-shrinking strategy quantization is one tool, not the only one.

💡 EXAMPLE

Running a 4-bit quantized model on a laptop = quantization. Combining quantization + pruning + distillation for a mobile model = full compression pipeline.

✖ COMMON MISTAKE

Saying "we compressed the model" when you only quantized it leaves out other available techniques.



Temperature **VS** Top-P

? WHAT

Temperature: scales the entire probability distribution to control randomness.

Top-P (nucleus sampling): truncates the distribution to the smallest set of tokens whose cumulative probability hits P.

✂ WHY THE CONFUSION

Both are "randomness knobs" for generation.

🕒 WHEN IT MATTERS

Tuning creative vs. deterministic output they interact, and tuning both without understanding each can give unpredictable results.

💡 EXAMPLE

Temperature 0.9 + Top-P 1.0 = broad, creative sampling. Temperature 0.9 + Top-P 0.3 = creative but restricted to only the most likely tokens.

✖ COMMON MISTAKE

Cranking both temperature and top-p high at once expecting more control usually just adds chaos.



AI Coach John

Follow



PROITBRIDGE
Strive For Better Future

 Repost

Structured Output **vs** JSON Mode

? WHAT

JSON Mode: guarantees the output is valid JSON but not any specific shape.

Structured Output: enforces a specific schema (defined fields, types, structure).

⚡ WHY THE CONFUSION

Both sound like "the model returns clean, parseable data."

🕒 WHEN IT MATTERS

Building production pipelines JSON mode alone can still give you wrong/missing fields; structured output enforces the exact contract your code expects.

💡 EXAMPLE

JSON mode might return `{"result": "yes"}` when you expected `{"status": "approved", "confidence": 0.9}`. Structured output forces that exact shape.

✖ COMMON MISTAKE

Relying on JSON mode alone for critical pipelines, then writing brittle code to handle unexpected shapes.



AI Coach John

Follow



 Repost

Long-Term Memory **VS** Context Window

? WHAT

Context Window: the temporary working memory for one session.

Long-Term Memory: a separate persistent store the system retrieves from and feeds back into context.

✂ WHY THE CONFUSION

Both feel like "what the model remembers."

🕒 WHEN IT MATTERS

Designing personalized AI products long-term memory needs its own storage + retrieval system; it doesn't just "grow" the context window.

💡 EXAMPLE

A coding assistant remembering your preferred code style across projects = long-term memory feeding facts back into a fresh context window each time.

✖ COMMON MISTAKE

Assuming long context window = long-term memory. Without a memory architecture, everything is forgotten once the session ends.



AI Coach John

Follow



 Repost