


GenAI Cheat Sheet

A Step-by-Step Learning Path for
Developers & Professionals



Foundation
Understand the core concepts

Communication
Master prompting techniques

Scale
Think like an engineer

Data Access
Leverage your data with RAG

Autonomy
Build agents that take action



PROITBRIDGE
Strive For Better Future



AI Coach John
Follow for more
Learn. Build. Ship. Repeat.

Understand LLMs

The Foundation Everything Else Builds On

HOW LLMs GENERATE TEXT



WHAT IS GENAI?

Generative AI is a class of AI systems that can generate new content — text, code, images, audio — by learning patterns from massive datasets.

Unlike traditional AI that classifies or predicts, GenAI creates.

Traditional AI

- classify
- predict
- analyze

Generative AI

- create
- generate
- synthesize


WHAT IS AN LLM?

A Large Language Model is a deep neural network trained on trillions of tokens of text.


It learns to predict the next token given a context.

At scale, this simple objective produces **emergent abilities**: reasoning, coding, translation, summarization, and more.

LLM CORE MECHANISM



TRANSFORMER MODEL (Simplified)



5 CORE CONCEPTS TO KNOW

1 TOKENS

The atomic unit of text an LLM reads and writes. Not words — subwords.




2 CONTEXT WINDOW

The maximum number of tokens the model can see at once — both input and output combined.



3 TEMPERATURE

Controls randomness of output.



Use low temperature for facts, high for creativity.

4 TRANSFORMER BASICS

The architecture behind every modern LLM.



5 TRAINING vs INFERENCE

Two very different phases.

TRAINING (Learning Phase)

- Trillions of tokens
- Weeks or months
- Thousands of GPUs
- Costs millions

INFERENCE (Usage Phase)

- Uses what it learned
- Milliseconds
- Fractions of a cent
- You interact here

MODELS TO KNOW

Model	Creator	Type	Highlight
GPT-4o	OpenAI	Proprietary	Multimodal
Claude 3.5	Anthropic	Proprietary	Strong reasoning
Gemini 1.5	Google	Proprietary	1M context
LLaMA 3	Meta	Open-source	Run locally
Mistral	Mistral AI	Open-source	Efficient
Qwen 2.5	Alibaba	Open-source	Multilingual

HOW MODELS GENERATE TEXT (Step-by-Step)






Prompt Engineering – The Most Underrated Developer Skill



THE CORE IDEA

The model's output quality is directly proportional to the quality of your input.



Prompt engineering is not a hack. It's a discipline.

CORE TECHNIQUES

1

ZERO-SHOT PROMPTING



Give the task directly with no examples. Works when the task is well-defined and the model has strong priors.



"Classify this email as spam or not spam: [email]"

2

FEW-SHOT PROMPTING



Provide 2–5 input-output examples before your real query. The model learns the pattern from examples.

Input: "I love this!" → Sentiment: Positive
Input: "This is awful" → Sentiment: Negative
Input: "It's okay" → Sentiment: ?

3

CHAIN-OF-THOUGHT (CoT)



Tell the model to think step by step before answering. Dramatically improves math, logic, and reasoning tasks.

"Solve this step by step: If a train travels 60 mph for 2.5 hours..."

4

ROLE PROMPTING



Assign the model a persona. Changes tone, depth, and framing of the response.

"You are a senior security engineer at a fintech company. Review this code for vulnerabilities."

5

STRUCTURED OUTPUT



Ask for a specific format — JSON, markdown table, numbered list, XML. Essential for programmatic use.

"Return your answer as a JSON object with keys: summary, sentiment, confidence_score"

6

PROMPT TEMPLATES



Build reusable prompt skeletons with variables. Separate the structure from the dynamic inputs.

"You are a {role}. Given the following {input_type}, {task}. Format: {format}."



BAD vs BETTER PROMPT



BAD PROMPT:

"Summarize this."

Problems:

- No audience defined.
- No length.
- No format.
- No tone.



Result: Random length, random style, may miss key points.



BETTER PROMPT:

"Summarize the following article in exactly 5 bullet points for a C-suite executive audience. Focus on business impact, risks, and recommended actions. Be direct and avoid jargon."

Why it works:

- Audience defined.
- Format specified.
- Focus specified.
- Tone specified.
- Length controlled.
- Result is predictable.



PROMPT ANATOMY



[ROLE]

→ Who the model is



[CONTEXT]

→ Background information



[INSTRUCTION]

→ The actual task



[FORMAT]

→ How to structure the output



[CONSTRAINTS]

→ What NOT to do



[EXAMPLES]

→ Optional few-shot demonstrations



Better prompts = better outputs.
You control quality before the model ever runs.





Production LLM Thinking

— Where Most Projects Break Down



THE HARD TRUTH



Demo AI

→ 1 user, 1 prompt,
works great in a notebook



Production AI

→ 10,000 users, 1,000 prompt variants,
cost constraints, latency SLAs,
failure modes

Most teams ship the demo and wonder why production is painful.
These are the concepts that separate engineers who build toy apps
from engineers who build **real systems**.

CORE CONCEPTS

1

PROMPT VERSIONING



Every prompt change is a code change.
Version your prompts in Git.
A/B test prompt versions.
Track which version produced
which output.
Never change a production prompt
without a rollback plan.

prompts/

```
├── summarize_v1.txt ← previous
├── summarize_v2.txt ← current (deployed)
└── summarize_v3.txt ← testing
```

2

CACHING



If 40% of your users ask the
same question, why call the
LLM 40% of the time?
Semantic caching matches
similar queries and returns
cached responses.
Saves cost and cuts latency
dramatically.



Exact cache:

"What is the capital of France?"
→ always Paris



Semantic cache:

"France capital?"
≈ "What's the capital of France?"

3

TOKEN COST OPTIMIZATION



Every token costs money.
Input tokens + Output
tokens = bill.

- Trim whitespace and redundant instructions from prompts
- Use cheaper models for simple tasks (GPT-4o-mini vs GPT-4o)
- Use streaming to improve perceived performance
- Set max_tokens to limit runaway output

4

LATENCY AWARENESS



LLMs are slow by software
standards. Time to First Token
(TTFT) feels like forever
at 3 seconds.

- Stream output token by token — feels 10× faster to the user
- Run non-dependent LLM calls in parallel
- Cache at multiple layers
- Use smaller models for latency-sensitive paths



Streaming makes 3s feel like 300ms

5

CONTEXT MANAGEMENT



Context windows are finite
and expensive. Sending
100K tokens every call
costs 25× more than 4K.

- Summarize old conversation turns instead of keeping all history
- Only inject the relevant retrieved chunks, not entire documents
- Use sliding window memory for long conversations



6

MONITORING



You cannot improve what
you cannot see.

Track:

- Input tokens
- Output tokens
- Latency (TTFT + Total)
- Error rate
- Cost per query
- Guardrail triggers
- User feedback signals

Measure → Understand → Improve



THE PRODUCTION CHECKLIST

- | | |
|---|--|
| ☐ Prompts versioned and tested before deploy | ☐ Retry logic with exponential backoff |
| ☐ Semantic cache in place for repeated queries | ☐ Fallback model configured (primary fails → secondary) |
| ☐ Streaming enabled for user-facing responses | ☐ Logging pipeline capturing inputs + outputs |
| ☐ Token budget set per request | ☐ Cost alerting set up |



GOAL

Build systems that scale,
cost predictably, and
don't break at 3am.




LLM Security: The Part Developers Skip Until It's Too Late



1 WHY THIS MATTERS NOW

Large Language Models introduce new security risks. Traditional validation methods are not enough. AI systems can be influenced through prompts, context, and user interactions.


Teams building GenAI products must design for security from the beginning.



2 SECURITY RISKS TO UNDERSTAND

1

PROMPT INJECTION



User instructions may attempt to override intended AI behavior and manipulate outputs.

Example:
"Ignore previous instructions" style prompt manipulation.

Defense:

- Input validation
- Output verification
- Permission separation

2

JAILBREAKING




Creative prompting techniques attempt to bypass model safeguards using indirect wording, fictional framing, or encoded instructions.

Defense:

- Safety classifiers
- Content moderation
- Guardrail systems

3

DATA LEAKAGE



AI systems may accidentally expose hidden instructions, retrieved context, or sensitive internal information.

Defense:

- Avoid storing secrets in prompts
- Use secure storage
- Limit exposed context

4

PII PROTECTION



Sensitive personal information may appear in prompts, outputs, or logs.

Defense:

- Detect and redact sensitive data
- Clean logs
- Use PII detection libraries

5

ACCESS CONTROL



Users should only access authorized information inside RAG pipelines and AI systems.

Defense:

- User scoped retrieval
- Permission filtering
- Organization level access checks

6

SAFETY GUARDRAILS




Additional validation layers run before and after model responses to reduce unsafe or unintended outputs.

Defense:

- Input filtering
- Output filtering
- Monitoring

3 DEFENSE IN DEPTH FLOW



4 KEY QUESTION

Can a malicious interaction **manipulate** your AI system, **expose** sensitive information, or **affect** other users?

If the answer is unclear, security testing should happen before deployment.



Hallucination & Evaluation – You Can't Trust What You Can't Measure

 **THE CORE PROBLEM**

Confident answer ≠ Correct answer
LLMs don't know what they don't know. They generate the most statistically likely next token — not the most factually accurate one.
A model can cite a paper that doesn't exist, give you a wrong API endpoint with complete confidence, or invent a legal case and quote it perfectly formatted.



WHAT IS HALLUCINATION?

Hallucination is when an LLM generates content that is factually incorrect, fabricated, or not supported by any provided context — but states it confidently as fact.

TYPES

**Factual Hallucination**
Wrong facts stated as truth.
"Einstein won the Nobel Prize for Relativity"

**Source Hallucination**
Cites references that don't exist.

**Context Hallucination**
Contradicts the document you provided (RAG failure).

**Instruction Hallucination**
Claims to have done something it didn't.

CORE EVALUATION CONCEPTS

1 GROUNDED RESPONSES



Every claim in the response should trace back to a source — the retrieved context, a document, or verified knowledge.

If you can't point to where a fact came from, it's a hallucination risk.


Retrieved Context

 →


Response

 →


Source Traceable

2 LLM EVALUATION (LLM-as-Judge)



Use a stronger model (GPT-4o, Claude) to evaluate another model's output.

Score on correctness, relevance, faithfulness, format compliance. Scalable, cheap, and surprisingly good.

Judge prompt:
"Rate the following response 1-5 for factual accuracy.
Reference answer: [gold].
Model response: [output].
Score + reason:"

3 HUMAN EVALUATION



The gold standard. Humans rate outputs on a rubric.

Expensive and slow, but required for calibrating automated evals and final validation before launch.

RUBRIC EXAMPLE

- 1 Poor
- 2 Fair
- 3 Good
- 4 Very Good
- 5 Excellent

4 ACCURACY CHECKS



Programmatic assertions on model output:

```
assert "disclaimer" in response.lower()
# must include disclaimer

assert len(response.split()) < 200
# must be concise

assert json.loads(response)["score"] >= 0
# must be valid JSON
```

5 BENCHMARKING



Standard test suites to compare models:


- MMLU — 57 academic subjects, tests breadth of knowledge
- HumanEval — Code generation correctness
- MT-Bench — Multi-turn conversation quality
- RAGAS — RAG-specific: faithfulness, context recall, answer relevance

6 EVALUATION DIMENSIONS

	Correctness	→	Is the answer factually right?
	Faithfulness	→	Is it grounded in provided context?
	Relevance	→	Does it answer the actual question?
	Completeness	→	Did it address all parts of the query?
	Conciseness	→	Is it free of unnecessary verbosity?
	Format	→	Did it follow the output format requested?
	Safety	→	Does it contain harmful or biased content?



RAG — Give Your LLM Access to Your Data



1 WHY RAG EXISTS

LLMs are frozen in time. They know nothing about your company's internal docs, your product, your customer data, or anything that happened after training. Fine-tuning is expensive and slow.

RAG solves this with retrieval at inference time — no retraining required.



Without RAG: LLM answers from training data only → stale, hallucinated



With RAG: LLM answers from your documents → grounded, current, accurate

2 THE SIMPLE RAG FLOW



3 CORE CONCEPTS

A EMBEDDINGS



Dense vector representations of text where semantic similarity = geometric proximity.

"King" and "Queen" are close.
"Dog" and "automobile" are far.

The embedding model converts text → a list of ~1500 numbers.

"How do I reset my password?"
→ [0.12, -0.87, 0.44, 0.91, ...]

"Steps to change my password"
→ [0.11, -0.85, 0.46, 0.89, ...]

→ These two vectors are very similar
→ retrieved together

B CHUNKING




Documents must be split into retrievable pieces.

Too large → low precision.

Too small → missing context.

Typical: 256-512 tokens with 50-100 token overlap between chunks.



Overlap preserves context

C VECTOR DATABASE



Stores embeddings and enables fast approximate nearest-neighbor search.

Can query millions of vectors in milliseconds.


Tool	Best For
Pinecone	Managed cloud, easy setup
Weaviate	Multi-modal, open source
pgvector	Already using PostgreSQL
Qdrant	Performance, open source
ChromaDB	Local development

D RETRIEVAL



Find the Top-K chunks whose embedding is closest to the query embedding.

Cosine similarity is the standard distance metric.



E CONTEXT INJECTION

Retrieved chunks are inserted into the LLM prompt as context:



System: "Answer using only the context below."
Context: [Chunk 1] ... [Chunk 2] ... [Chunk 3]
User: "What is the refund policy?"



LLM Generates
Grounded Answer



GOAL

Connect your LLM to your data.
Ground every response. Eliminate guessing.





Advanced RAG — Beyond Simple Vector Search

THE UNCOMFORTABLE TRUTH

Naive RAG (embed → retrieve → inject) breaks in production.

Retrieval returns wrong chunks. Queries are ambiguous. Rankings are off. Metadata is ignored. The LLM gets bad context and hallucinates anyway — just with citations.



Basic RAG:
50–60%
retrieval accuracy
(typical naive implementation)




Advanced RAG:
80–90%+
with the right enhancements

ADVANCED CONCEPTS

1 HYBRID SEARCH

Combine dense vector search (semantic similarity) with sparse keyword search (BM25). Neither alone is best.




2 RE-RANKING

After retrieving Top-20 chunks by vector similarity, run a cross-encoder reranker to re-score them with full attention over (query, chunk) pairs. The cross-encoder is slower but far more accurate than embedding similarity. Return the re-ranked Top-5 to the LLM.




3 QUERY EXPANSION

The user's query is often a bad search query. Rewrite or expand it before retrieval.



4 METADATA FILTERING

Vector search alone ignores structure. Metadata filters add hard constraints before or after retrieval.



5 EVALUATION FOR RAG

Standard LLM eval isn't enough. RAG needs component-level evaluation:


Metric	What it Measures
Context Precision	Are retrieved chunks actually relevant?
Context Recall	Were all relevant chunks retrieved?
Faithfulness	Does the answer only use retrieved context?
Answer Relevance	Does the answer address the question?

Use RAGAS framework to measure all four automatically.

6 ADVANCED RAG PATTERNS

- HyDE**: Generate a hypothetical answer → embed it → retrieve similar real docs
- Parent-Child**: Retrieve small chunks → return their parent chunk for full context
- Step-Back**: Ask a broader question first → retrieve general then specific
- Multi-Query**: Generate 3 versions of the query → retrieve for all → merge
- FLARE**: LLM generates → pauses when uncertain → retrieves → continues
- GraphRAG**: Build knowledge graph → retrieve via entity relationships

HOW IT ALL WORKS TOGETHER



GOAL
Good RAG is an engineering discipline — not a one-liner. Build it right.



AI Agents — When the LLM Takes the Wheel



1 WHAT IS AN AGENT?

An agent is an LLM that doesn't just answer — it acts. It can decide what tool to use, use it, observe the result, decide what to do next, and repeat until the goal is achieved. The LLM is the brain. Tools are the hands.

Single prompt:
User asks → LLM answers → Done

Agent:
User sets goal → LLM plans → executes tools → observes → re-plans → executes → ... → delivers result




2 CORE CONCEPTS

2.1 TOOL CALLING (FUNCTION CALLING)

The LLM can call external functions — search the web, query a database, run Python, call an API, read a file. The model outputs a structured JSON tool call. Your code executes it. The result comes back. The model continues.

Model output:

```
{  "tool": "search_web",  "arguments": {    "query": "current NVIDIA stock price"  }}
```



2.2 MEMORY

Agents need to remember. Four types:

- In-context:** The current conversation — limited by context window
- External (Vector):** Long-term retrieval from a vector DB across sessions
- Entity Memory:** Track specific facts (user name, preferences, past decisions)
- Episodic:** Log of what the agent did before (actions + outcomes)

2.3 PLANNING

Breaking a complex goal into executable sub-tasks.

Goal: "Write a research report on LLM safety"
Plan:

1. Search recent papers on LLM safety (tool: search_web)
2. Read and summarize each paper (tool: read_url)
3. Extract key themes across papers (LLM reasoning)
4. Draft report structure (LLM reasoning)
5. Write each section (LLM generation)
6. Save final report (tool: write_file)


2.4 MULTI-STEP REASONING

Agents use the ReAct loop → Reason → Act → Observe → Repeat.



2.5 WORKFLOW AUTOMATION

Agents can replace multi-step manual workflows:



2.6 EVALUATION FOR AGENTS

Agents need evaluation too. Track what matters.

Metric	What it Measures
 Task Success Rate	Did the agent complete the task correctly?
 Tool Selection Accuracy	Did the agent choose the right tools?
 Step Efficiency	Did the agent take the optimal number of steps?
 Hallucination Rate	Did the agent generate incorrect information?
 Cost per Task	Total tokens + tool costs per completed task
 User Satisfaction	Was the user happy with the result?

3 AGENT EXAMPLES

Agent	What it Does	Tools Used	Example Use Case
 Research Agent	Searches, reads, synthesizes information from multiple sources	web_search, read_url, write_file, vector_db, summarizer	Build market research reports, literature reviews
 Coding Agent	Writes, runs, debugs, fixes, and refactors code	code_interpreter, file_system, linter, git, debugger	Build full-stack features, fix bugs, refactor codebase
 Customer Support Agent	Reads tickets, queries knowledge, drafts replies, handles escalations	CRM API, email, knowledge base, search, templates	Answer customer queries, route issues, update CRM
 Data Agent	Pulls data, analyzes, creates visualizations and insights	SQL, Python, pandas, plotting libraries, data APIs	Sales reports, KPI dashboards, trend analysis
 DevOps Agent	Monitors infrastructure, detects issues, takes corrective action	cloud APIs, monitoring tools, alerting, runbooks, shell	Incident response, auto-scaling, alert triage, deployments



GOAL

Move beyond single prompts.
Build systems that take initiative, use tools, and deliver results.



Production Agent Design — Powerful Agents Need Strong Controls



⚠ THE PRODUCTION REALITY

A demo agent that sometimes works is impressive. A production agent that reliably works for 10,000 users under adversarial conditions — that's engineering.

Single-agent demos are easy. Production multi-agent systems with real users, real data, and real consequences are an entirely different challenge.



1 MULTI-AGENT SYSTEMS

Decompose complex goals across specialist agents that collaborate. Each agent has a narrow role and defined interface.

Orchestrator Agent
(plans, delegates, synthesizes)

Research Agent

Analysis Agent

Writer Agent

Review Agent

Frameworks: LangGraph · CrewAI · AutoGen · AgentTorch

2 HUMAN-IN-THE-LOOP (HITL)

Not all decisions should be autonomous. High-stakes actions need human approval. Design explicit checkpoints where the agent pauses and asks.

Low risk: Agent sends a Slack message	No approval needed
Medium risk: Agent creates a Jira ticket	Approval optional
High risk: Agent deploys code to production	Requires human sign-off

Agent proposes an action

Human reviews & approves

Agent continues execution

3 GUARDRAILS

Agents can take real actions with real consequences — sending emails, deleting files, charging customers. Guardrails define what the agent is allowed to do.

Allowed:
read files from /data/reports

Blocked:
delete any file

Blocked:
send emails to external domains

Blocked:
spend more than \$50 per task without approval

Tools:
LangSmith · Langfuse · Arize · Helicone

4 OBSERVABILITY

You must be able to see exactly what every agent did, in what order, with what inputs and outputs.

Trace every agent call:

- Which tool was called?
- What were the arguments?
- What was the result?
- How long did it take?
- What did it cost?
- Did it hit a guardrail?

5 FAILURE HANDLING

Agents fail. Tools time out. APIs return errors. LLMs hallucinate a tool argument. Design for failure explicitly.

Failure patterns to design for:

- Tool call returns error → retry with backoff → fallback tool → ask human
- Agent loops without progress → step limit → timeout → alert operator
- LLM produces invalid JSON for tool call → parse error → retry with correction prompt
- Goal is ambiguous → agent asks for clarification before acting

📋 PRODUCTION AGENT CHECKLIST

- 🔧 Every tool has input validation and error handling
- ∞ Maximum step limit set (prevent infinite loops)
- 💰 Maximum cost per task limit enforced
- 👤 High-risk actions require human approval
- 📄 Full trace logging for every agent run
- 🔒 Agent cannot access data beyond its permission scope
- 🔄 Fallback path defined for every critical tool
- 📊 Eval suite running on agent traces post-deploy
- 🔔 Alerting on failure rate, cost spikes, loop detection
- 🔄 Rollback plan if a bad agent run causes damage

✅ THE FULL GENAI LEARNING PATH — COMPLETED

✓	Step 1	→	LLM Fundamentals	Foundation	🏛️
✓	Step 2	→	Prompt Engineering	Communication	💬
✓	Step 3	→	Production Thinking	Scale	📊
✓	Step 4	→	Security & Safety	Protection	🛡️
✓	Step 5	→	Hallucination & Eval	Quality	🎯
✓	Step 6	→	RAG Fundamentals	Data Access	🗄️
✓	Step 7	→	Advanced RAG	Better Retrieval	🔍
✓	Step 8	→	AI Agents	Autonomy	🤖
✓	Step 9	→	Production Agents	Reliability	🛡️



FINAL REALITY CHECK

Anyone can prompt an LLM.

Few can build systems that are secure, reliable, and production-grade.

You now have the map. The rest is execution.



Powerful agents need strong controls.

Build responsibly.

