


LangGraph Core Concepts: The Complete Guide

Every primitive you need to read and build state-of-the-art
LangGraph AI applications.



 LangGraph



AI Coach John



FOLLOW

 Repost

Every Concept LangGraph Uses

The 12 foundational primitives of stateful agentic graphs



1. State

Shared memory dict passed between steps.



2. Nodes

Functions that execute work and update state.



3. Edges

Fixed deterministic execution paths.



4. Conditional

Dynamic runtime branching based on state.



5. Reducers

Rules for merging concurrent state writes.



6. Parallel

Fan-out/fan-in dispatch via Send API.



7. Checkpoint

Thread memory persistence across turns.



8. Streaming

Intermediate event emitting per step.



9. Interrupts

Human-in-the-loop pause & resume.



10. ToolNode

Prebuilt LLM tool execution node.



11. Limits

Recursion safety caps against loops.



12. Subgraphs

Compiled graphs nested inside nodes.



Core Blueprint: Master these 12 building blocks to architect any cyclical, multi-agent, or human-supervised AI workflow.



AI Coach John



FOLLOW

 Repost

State: Shared Memory

The single schema passed between every node

What It Is

A single data structure (usually a TypedDict) passed between every node. Each node reads from it and writes back to it: the only way information flows through the graph.

```
from typing import TypedDict

class GraphState(TypedDict):
    question: str
    retrieved_docs: str
    answer: str
```

State (in)
{question: "..."} →

Node runs
reads + updates dict →

State (out)
{question, answer}

Why Explicit Schemas Matter

Every field a node might need should be declared explicitly in GraphState. Avoiding global variables ensures clean state isolation and predictable testing.



Architecture Rule: Design your state schema before writing any nodes. Explicit fields prevent hidden state bugs.



AI Coach John



FOLLOW

 Repost

Nodes: Execution Steps

Individual functions that perform work

What It Is

A node is a Python function that takes the state, performs an action (calls an LLM, queries a tool, runs logic), and returns an updated state.

```
def retrieve_node(state: GraphState) -> GraphState:
    docs = retriever.get_relevant_documents(state["question"])
    state["retrieved_docs"] = "\n".join([d.page_content for d in docs])
    return state
```

```
graph.add_node("retrieve", retrieve_node)
graph.add_node("generate", generate_node)
```

"retrieve" node
fetches context



"generate" node
synthesizes answer

Single-Responsibility Rule

Resist cramming retrieval, grading, and generation into one monolithic node. Keeping nodes small enables individual unit testing and clean graph loops.



Best Practice: Each node should do one focused thing. Modularity is what separates LangGraph from spaghetti chains.



AI Coach John



FOLLOW

 **Repost**

Edges & START

Fixed deterministic connections between nodes

What It Is

An edge says: "after this node finishes, always run that node next." START and END are built-in constants marking entry and exit points.

```
from langgraph.graph import StateGraph, START, END

graph = StateGraph(GraphState)
graph.add_node("retrieve", retrieve_node)
graph.add_node("generate", generate_node)

graph.add_edge(START, "retrieve") # explicit entry point
graph.add_edge("retrieve", "generate")
graph.add_edge("generate", END)
```

START



retrieve



generate



END

Deterministic Topology

Using explicit START edges instead of legacy entry methods makes your graph diagram immediately readable and ready for runtime branching.



Modern Standard: Explicit START and END edges make your graph topology visually clear at a glance.



AI Coach John



FOLLOW

 Repost

Conditional Edges

Dynamic runtime branching based on state inspection

What It Is

Instead of always going to the same next node, a conditional edge evaluates a routing function against state and selects the destination dynamically.

```
def route_after_grading(state: GraphState) -> str:
    return "generate" if state["sufficient"] else "retrieve"

graph.add_conditional_edges(
    "grade",
    route_after_grading,
    {"generate": "generate", "retrieve": "retrieve"}
)
```

retrieve



grade



generate
(if sufficient)

 Insufficient context: loop back from grade → retrieve



Runtime Guardrail: Every possible return string from your routing function must exist in the edge mapping dictionary to prevent runtime `KeyErrors`.



AI Coach John



FOLLOW

 **Repost**

Reducers: State Merging

Controlling how concurrent writes get combined

What It Is

When multiple nodes write to the same state key, LangGraph needs a reducer function to define whether to overwrite or combine those writes.

```
import operator
from typing import Annotated, TypedDict

class GraphState(TypedDict):
    question: str
    # operator.add merges list writes by concatenation
    sources: Annotated[list[str], operator.add]
```

Node A writes:

`sources = ["doc1"]`



Node B writes:

`sources = ["doc2"]`



Reducer (operator.add)

`sources = ["doc1", "doc2"]`

Preventing Silent Overwrites

Without an explicit reducer annotation, LangGraph defaults to "last write wins." Annotating fields with reducers like `operator.add` safely accumulates results.



Essential Primitive: Reducers are mandatory whenever you run parallel branches or append conversational message histories.



AI Coach John



FOLLOW

 **Repost**

Parallel Execution: Send API

Fan-out / fan-in dispatch for concurrent processing

What It Is

The Send API lets one node dispatch multiple concurrent instances of another node, each handling a different slice of state, then merges outputs.

```
from langgraph.constants import Send

def dispatch_subquestions(state: GraphState):
    return [Send("answer_subquestion", {"sub_q": q}) for q in
state["sub_questions"]]

graph.add_conditional_edges("decompose", dispatch_subquestions,
["answer_subquestion"])
```

decompose



Fan-Out: sub-Q 1, 2, 3
parallel node instances



combine

Latency Optimization

Instead of processing 4 sub-queries sequentially (taking 12s total), parallel Send dispatch runs all 4 concurrently, completing in just ~3s.



Speed Advantage: Essential for multi-hop RAG decomposition. Runs sub-questions concurrently instead of sequentially.



AI Coach John



FOLLOW

 Repost

Checkpointing: Thread Memory

Persisting graph state across turns and sessions

What It Is

A checkpointer saves state after every step keyed by `thread_id`, enabling pause/resume, human review, and multi-turn conversational memory.

```
from langgraph.checkpoint.memory import MemorySaver

checkpointer = MemorySaver()
app = graph.compile(checkpointer=checkpointer)

config = {"configurable": {"thread_id": "user-123"}}
res1 = app.invoke({"question": "Refund policy?"}, config=config)
res2 = app.invoke({"question": "What about digital?"}, config=config)
```

Invocation 1
retrieve → generate



Saved Checkpoint
thread_id: "user-123"



Invocation 2
resumes with full memory

Production Checkpointers

While `MemorySaver` is ideal for local testing, production deployments should use persistent database checkpointers (Postgres, SQLite, or Redis).



Production Note: Checkpointing is what transforms stateless functions into persistent multi-turn assistants.



AI Coach John



FOLLOW

 **Repost**

Streaming: Real-Time Progress

Surface node updates as the graph runs

What It Is

.stream() yields output after each node completes, rather than waiting for the entire multi-step loop to finish before returning.

```
for event in app.stream({"question": "Compare Q1 & Q3"}, config=config):  
    for node_name, output in event.items():  
        print(f"[{node_name}] → {output}")
```

retrieve
"Searching..."



grade
"Checking..."



generate
"Answer ready"

Transparent User Experience

Long-running agent loops can take 5 to 15 seconds. Streaming intermediate node statuses prevents perceived freezes and builds user trust.



UX Imperative: Surface intermediate node updates so users see real-time reasoning rather than a spinning loader.



AI Coach John



FOLLOW

 Repost

Interrupts: Human Review

Pausing execution for approval before side effects

👤 What It Is

`interrupt_before` pauses execution at a named node, letting a human inspect or modify state before resuming execution from checkpoint.

```
app = graph.compile(checkpointer=checkpointer, interrupt_before=[
    "send_email"
])
```

```
# Graph pauses right before "send_email":
app.invoke({"question": "Draft refund email"}, config=config)
```

```
# Inspect state, approve, then resume:
app.invoke(None, config=config)
```

draft_email



PAUSED (Human Review)

inspect & approve state



send_email



Safe Resume Workflow: Calling `app.invoke(None, config)` resumes execution exactly where it paused without re-running earlier nodes.



AI Coach John



FOLLOW

Repost

ToolNode: Prebuilt Execution

Standardized LLM tool calling and message wiring

What It Is

Instead of hand-writing custom parsing logic, ToolNode automatically executes tool calls requested by the LLM and appends results.

```
from langgraph.prebuilt import ToolNode

tools = [retrieve_docs, web_search, sql_query_tool]
tool_node = ToolNode(tools)

graph.add_node("tools", tool_node)
graph.add_conditional_edges("agent", lambda s: "tools" if s["messages"]
[-1].tool_calls else END)
graph.add_edge("tools", "agent")
```

agent
(LLM w/ tools)



ToolNode
executes tool_calls



END



Automatic Formatting: ToolNode executes tools and automatically formats outputs as valid ToolMessage objects for OpenAI & Anthropic.



AI Coach John



FOLLOW

 **Repost**

Recursion Limit: Safety Cap

Hard backstop against infinite agent loops

What It Is

A hard limit capping the maximum number of steps a graph can execute in a single run, preventing runaway billing from cyclical loops.

```
result = app.invoke(  
    {"question": "..."},  
    config={"configurable": {"thread_id": "u-1"}, "recursion_limit": 25}  
)  
# raises GraphRecursionError if exceeded
```

Why Retry Counters Aren't Enough

Even if you track retries in your state dictionary, a bug in conditional routing logic can trap the graph in an infinite cycle. Explicit recursion caps protect production stability.

Tuning Your Cap

For standard RAG pipelines, set `recursion_limit: 10`. For multi-agent autonomous reasoning loops, allow 25 to 50 steps.



Production Guardrail: Always set an explicit `recursion_limit` tailored to your expected maximum graph depth.



Subgraphs: Modular Nodes

Nesting compiled graphs inside parent graphs

What It Is

A fully compiled graph can be added as a node inside a larger parent graph. This is the foundation of multi-agent supervisor architectures.

```
docs_agent = docs_graph.compile()    # compiled independent graph
sql_agent = sql_graph.compile()

supervisor_graph = StateGraph(SupervisorState)
supervisor_graph.add_node("docs_agent", docs_agent)
supervisor_graph.add_node("sql_agent", sql_agent)
```

Supervisor
parent graph



docs_agent (Subgraph)
isolated state & schema



sql_agent (Subgraph)
isolated state & schema

Team & Domain Isolation

Subgraphs allow different engineering teams to independently build, test, and version their domain graphs before integrating them into a master supervisor.



Encapsulation: Subgraphs maintain their own schema, checkpointers, and internal loops, keeping enterprise systems maintainable.



AI Coach John



FOLLOW

 Repost

Putting It All Together

A production graph combining all core concepts

```
class GraphState(TypedDict):
    question: str; sub_questions: list[str]; answer: str
    sources: Annotated[list[str], operator.add] # reducer

graph = StateGraph(GraphState)
graph.add_node("decompose", decompose_node)
graph.add_node("answer_subq", subq_node) # parallel via Send
graph.add_node("grade", grade_node)
graph.add_node("generate", generate_node)
graph.add_node("tools", ToolNode(tools))

graph.add_edge(START, "decompose")
graph.add_conditional_edges("decompose", dispatch_subq, ["answer_subq"])
graph.add_edge("answer_subq", "grade")
graph.add_conditional_edges("grade", route_after_grading, {"generate": "generate",
"decompose": "decompose"})
graph.add_edge("generate", END)

app = graph.compile(checkpointer=MemorySaver(), interrupt_before=["generate"])
```

✓ Architectural Synergy

State + reducers handle concurrent fan-out writes; nodes execute actions; conditional routing self-corrects; checkpointers + interrupts add human oversight.



Complete System: Ready for real-time streaming, pause/resume, and safe production execution.



AI Coach John



FOLLOW

Repost

Quick Reference Matrix

All 12 primitives at a glance

CONCEPT	WHAT IT DOES	KEY API / METHOD
State	Shared memory between nodes	TypedDict class
Nodes	Functions performing work	add_node(name, fn)
Edges / START	Fixed deterministic paths	add_edge(from, to)
Conditional	Dynamic state-based routing	add_conditional_edges()
Reducers	Merge rule for parallel writes	Annotated[type, reducer]
Parallel Execution	Fan-out concurrent runs	Send(node, state)
Checkpointing	Persist thread state	compile(checkpointer=...)
Streaming	Intermediate node progress	app.stream(...)
Interrupts	Human review pause	interrupt_before=[...]
ToolNode	Prebuilt LLM tool runner	ToolNode(tools)
Recursion Limit	Loop safety backstop	recursion_limit config
Subgraphs	Nested compiled graph	add_node("x", graph)



Developer Reference: Keep this matrix handy when debugging graph syntax or designing state schemas.



AI Coach John



FOLLOW

Repost

 WATCH OUT

Common LangGraph Pitfalls

Mistakes that trip up developers early

Missing Return State

Forgetting to return updated state dict from a node leaves state unchanged.

Missing Reducer

No reducer on parallel fields causes silent "last write wins" data loss.

Unmapped Router Key

Conditional edge function returning an unmapped key crashes at runtime.

No Recursion Cap

No retry cap + no recursion limit can trap graphs in endless loop billing.

In-Memory Prod Storage & Monolithic Nodes

Using `MemorySaver` in production loses state on restart. Overloading single nodes defeats modularity. Use `ToolNode` instead of hand-rolling tool logic.



Action: Save this guide as your complete LangGraph reference. Follow for more AI engineering deep dives!



AI Coach John



FOLLOW

 Repost