



AI ENGINEERING STACK

Build a Production RAG Pipeline with LangChain

Step-by-step from raw documents to intelligent answers

Swipe to build it yourself →

AI Coach John

Follow

PROITBRIDGE

↻ Repost

What is RAG?

Retrieval-Augmented Generation. Instead of relying on an LLM's training data alone, you fetch relevant documents at runtime and feed them as context. The LLM answers from your data, not from memory.

★ WHY IT MATTERS

- ✓ No hallucinations from stale training
- ✓ Works on private/custom data
- ✓ Always up to date

Pipeline preview:

Document → Chunk → Embed → Store
Query → Retrieve → Validate → Generate

Install Dependencies

Get everything you need in two commands. We use HuggingFace for free embeddings, FAISS for vector storage, and Groq for a fast free LLM.

```
pip install langchain langchain_community  
pip install langchain_huggingface faiss-cpu  
pip install langchain_groq pypdf presidio-analyzer  
pip install presidio-anonymizer rank_bm25
```

STEP 2

Load Your Documents

Document loaders convert files into LangChain `Document` objects with `page_content` and `metadata`. Pick the loader that matches your source.

```
from langchain_community.document_loaders import PyPDFLoader

loader = PyPDFLoader("your_file.pdf")
documents = loader.load()

# Each document has:
# doc.page_content → the text
# doc.metadata      → {"source": "file.pdf", "page": 0}
print(f"Loaded {len(documents)} pages")
```



Other loaders: `TextLoader`, `WebBaseLoader`, `CSVLoader`, `WikipediaLoader`

Preprocess & Clean

LangChain has no built-in cleaner - write your own. Remove noise like page numbers, headers, and extra whitespace before chunking.

Garbage in = garbage retrieved.

```
import re

def clean(text):
    text = re.sub(r'\s+', ' ', text)      # extra whitespace
    text = re.sub(r'Page \d+', '', text)  # page numbers
    text = re.sub(r'^\w\s\.,\!\?]', '', text)
    return text.strip()

for doc in documents:
    doc.page_content = clean(doc.page_content)
```

Validate & Mask PII

Before anything touches the pipeline, sanitize the user query. Block prompt injections and mask personal data (emails, phones, IDs) to stay GDPR/HIPAA compliant.

```
from presidio_analyzer import AnalyzerEngine
from presidio_anonymizer import AnonymizerEngine

analyzer = AnalyzerEngine()
anonymizer = AnonymizerEngine()

def mask_pii(text):
    results = analyzer.analyze(text=text, language="en")
    return anonymizer.anonymize(text=text,
                                analyzer_results=results).text

def validate_query(query):
    banned = ["ignore previous", "jailbreak", "disregard"]
    for phrase in banned:
        if phrase in query.lower():
            raise ValueError("Invalid query detected")
    return mask_pii(query.strip())

safe_query = validate_query(user_query)
```

✂ STEP 5

Chunk the Documents

Split cleaned documents into smaller overlapping pieces. Chunk size controls precision; overlap preserves context across boundaries. Start at 300 chars and tune from there.

```
from langchain.text_splitter import
RecursiveCharacterTextSplitter

splitter = RecursiveCharacterTextSplitter(
    chunk_size=300,
    chunk_overlap=50,
    separators=["\n\n", "\n", ".", " ", " "]
)

chunks = splitter.split_documents(documents)
print(f"{len(chunks)} chunks created")
```

💡 **Rule of thumb:** `chunk_overlap` = 15–20% of `chunk_size`

Add Metadata to Chunks

Tag every chunk with structured metadata - source, category, date, author. This powers filtered retrieval later so you only search relevant subsets.

```
for i, chunk in enumerate(chunks):
    chunk.metadata.update({
        "chunk_id": i,
        "category": "langchain",    # tag manually or extract
        "year": 2024,
        "source": "rag_guide.pdf"
    })

# Now each chunk carries:
# {"source": "rag_guide.pdf", "category": "langchain",
#  "year": 2024, "page": 0, "chunk_id": 42}
```

STEP 7

Create Embeddings

Convert each chunk into a numerical vector that captures its meaning. Similar chunks end up close together in vector space - that's what enables semantic search.

```
from langchain_huggingface import HuggingFaceEmbeddings

embeddings = HuggingFaceEmbeddings(
    model_name="sentence-transformers/all-mpnet-base-v2",
    model_kwargs={"device": "cpu"}
)

# Test it
vector = embeddings.embed_query("What is RAG?")
print(f"Vector size: {len(vector)}") # 768 dimensions
```



Free alternatives: [all-MiniLM-L6-v2](#), [text-embedding-ada-002](#)

Store in Vector DB

Index all chunk embeddings into FAISS for fast similarity search. In production, swap FAISS for Chroma or Qdrant which support richer metadata filtering.

```
from langchain_community.vectorstores import FAISS

# Build the index from chunks + embeddings
vectorstore = FAISS.from_documents(chunks, embeddings)

# Save to disk so you don't rebuild every run
vectorstore.save_local("faiss_index")

# Load later
vectorstore = FAISS.load_local("faiss_index", embeddings,
                               allow_dangerous_deserialization=True)
```

Hybrid Search (BM25 + Dense)

Pure vector search misses exact keyword matches. BM25 catches those. Combine both with an ensemble retriever - you get the best of semantic meaning AND keyword precision.

```
from langchain_community.retrievers import BM25Retriever
from langchain.retrievers import EnsembleRetriever

bm25 = BM25Retriever.from_documents(chunks)
bm25.k = 5

dense = vectorstore.as_retriever(search_kwargs={"k": 5})

retriever = EnsembleRetriever(
    retrievers=[bm25, dense],
    weights=[0.4, 0.6] # tune based on your data
)
```



Tip: increase BM25 weight for technical/exact queries, dense for conceptual ones.

▼ STEP 10

Metadata Filtering

Narrow retrieval to a specific subset before similarity search runs. This is especially powerful for multi-tenant apps or domain-separated corpora.

```
# Use Chroma for production metadata filtering
from langchain_community.vectorstores import Chroma

vectorstore = Chroma.from_documents(chunks, embeddings)

filtered_retriever = vectorstore.as_retriever(
    search_kwargs={
        "k": 5,
        "filter": {
            "$and": [
                {"category": {"$eq": "langchain"}},
                {"year": {"$gte": 2023}}
            ]
        }
    }
)
```

</> Operators: `$eq`, `$ne`, `$gt`, `$gte`, `$lt`, `$lte`, `$in`, `$and`, `$or`

Multi-Query Retrieval

A single query often misses relevant chunks phrased differently. Generate 3–5 rephrasings automatically, retrieve for each, then deduplicate – recall goes up significantly.

```
from langchain.retrievers.multi_query import MultiQueryRetriever

multi_retriever = MultiQueryRetriever.from_llm(
    retriever=retriever,
    llm=llm
)

# Query: "How does RAG work?"
# Auto-generates:
# → "Explain the RAG architecture"
# → "What is retrieval augmented generation?"
# → "Steps in a RAG pipeline"
# Retrieves all, deduplicates, returns combined set

docs = multi_retriever.invoke(safe_query)
```

Contextual Compression

Retrieved chunks often contain irrelevant sentences alongside the useful ones. Compression strips those out - only the sentences that answer the query pass through to the LLM.

```
from langchain.retrievers import ContextualCompressionRetriever
from langchain.retrievers.document_compressors import
EmbeddingsFilter

# No LLM call needed - uses embedding similarity
compressor = EmbeddingsFilter(
    embeddings=embeddings,
    similarity_threshold=0.76
)

compression_retriever = ContextualCompressionRetriever(
    base_compressor=compressor,
    base_retriever=retriever
)

compressed_docs = compression_retriever.invoke(safe_query)
```

⚡ **Benefit: fewer tokens sent to LLM → lower cost + sharper context.**

Set Up the LLM

Initialize your LLM with temperature=0 for deterministic, factual answers. Groq gives you fast free inference on open models - great for prototyping.

```
from langchain_groq import ChatGroq

llm = ChatGroq(
    model_name="gemma2-9b-it",
    temperature=0,          # 0 = factual, 1 = creative
    api_key="your_groq_api_key"
)

# Test
response = llm.invoke("What is LangChain?")
print(response.content)
```



Other options: [ChatOpenAI](#), [ChatAnthropic](#), [ChatOllama](#)

Generate the Answer

Wire the retriever and LLM together with RetrievalQA. The chain fetches relevant chunks, injects them into the prompt, and the LLM answers from that context only.

```
from langchain.chains import RetrievalQA

rag_chain = RetrievalQA.from_chain_type(
    llm=llm,
    chain_type="stuff",          # all chunks in one prompt
    retriever=compression_retriever,
    return_source_documents=True
)

result = rag_chain.invoke({"query": safe_query})

answer = result["result"]
sources = result["source_documents"]
```

 Chain types: [stuff](#) (fast), [map_reduce](#) (long docs), [refine](#)

Validate the Output

Don't return LLM responses blindly. Check faithfulness against retrieved context and strip any PII that might have leaked through from source documents.

```
from langchain_core.prompts import PromptTemplate
from langchain_core.output_parsers import StrOutputParser

check_prompt = PromptTemplate.from_template("""
Context: {context}
Answer: {answer}
Is this answer supported by the context?
Reply ONLY: SUPPORTED or NOT_SUPPORTED
""")

checker = check_prompt | llm | StrOutputParser()

def validate_output(answer, docs):
    context = " ".join([d.page_content for d in docs])
    verdict = checker.invoke({"context": context, "answer": answer})
    if "NOT_SUPPORTED" in verdict:
        return "Could not find a reliable answer in the documents."
    return mask_pii(answer) # mask any leaked PII

final_answer = validate_output(answer, sources)
```

Full Pipeline at a Glance

Every card in this carousel is one step. Put them together and you have a production-ready RAG system.

1. Install dependencies
2. Load documents
3. Preprocess & clean
4. Validate input + mask PII ← security
5. Chunk documents
6. Tag chunks with metadata
7. Create embeddings
8. Store in vector DB
9. Hybrid search (BM25 + dense) ← recall
10. Metadata filtering ← precision
11. Multi-query retrieval ← coverage
12. Contextual compression ← efficiency
13. Set up LLM
14. Generate answer
15. Validate output + mask PII ← trust



Save this. Share it. Build with it.

Follow for more AI engineering content